

Natural-Language Mathematical Statements to Formal Statements

Autoformalization as alignment, not translation

XMUM GE – Week 8

Ma Jiajun – May 26, 2026

Part I. Opening Case Study

One symbol, many formal objects

One Symbol, Many Formal Objects: Integral

On the board only this:

$$\int_a^b f(x) dx$$

Question: to write this as a Lean theorem statement, what is missing?


The missing piece is not Lean syntax. It is the mathematical object.

Possible readings

1. Riemann integral on a compact interval.
2. Lebesgue / Bochner integral wrt a measure.
3. Oriented interval integral $\int_a^b$.
4. Set integral over $[a, b]$, $(a, b]$, ...
5. Variable upper limit $x \mapsto \int_a^x f(t) dt$.
6. Antiderivative relation $F' = f$.


Riemann vs Lebesgue: Two Ways To Sum The Area

Riemann – slice the x -axis



$$\int_a^b f dx \approx \sum_i f(x_i) \Delta x_i$$

Lebesgue – slice the y -axis



$$\int_E f d\mu = \int_0^\infty \mu(\{f > t\}) dt$$

Riemann: partition the **domain**. **Lebesgue:** partition the **range**; integrate $f \geq 0$ via the layer-cake formula above; for sign-changing f split $f = f^+ - f^-$ and require $\int |f| d\mu < \infty$. Conditionally convergent series fail this – Lebesgue refuses; improper Riemann still accepts.

Mathlib Says "As $R \rightarrow \infty$ " With Filter and Tendsto

Natural sentence

$$\int_1^R \frac{\sin x}{x} dx \rightarrow L \quad \text{as } R \rightarrow \infty$$

Mathlib writes

```
Tendsto (fun R => ∫ x in 1..R, sin x / x) atTop (ℵ L)
```

A function, a *source filter*, a *target filter*. No ε - δ on the surface.

Filter = "eventually"

- `atTop : Filter ℝ` – R is arbitrarily large
- `ℵ L : Filter ℝ` – value lies in any neighborhood of L

`Tendsto f F G` reads: *f* pushes "eventually in F " to "eventually in G ".

Same shape, different filters: `Tendsto u atTop atTop` = sequence to ∞ . `Tendsto f (ℵ a) (ℵ b) = f` continuous at a with $f(a) = b$. ε - δ lives inside the filters; the *statement* uses filter algebra.

Running Example: Improper Riemann Exists, Lebesgue Does Not

$$\int_1^{\infty} \frac{\sin x}{x} dx \text{ exists.}$$

Reading A – improper Riemann integral: **true**

```
∃ L : ℝ,  
  Tendsto  
    (fun R : ℝ ⇒ ∫ x in (1:ℝ)..R, Real.sin x / x)  
    atTop  
    (ℳ L)
```

Finite interval integrals have a limit as $R \rightarrow \infty$.

Reading B – Lebesgue integrable on $[1, \infty)$: **false**

```
¬ IntegrableOn  
  (fun x : ℝ ⇒ Real.sin x / x)  
  (Set.Ici 1)
```

$\int_1^{\infty} |\sin x/x| dx = \infty$. Lebesgue integrability needs absolute integrability; this is only conditional convergence.

Putnam 2025 B2: Geometry Becomes Integral Ratios

$f : [0, 1] \rightarrow [0, \infty)$ strictly increasing continuous.

R = region under $y = f(x)$. $x_1, x_2 = x$ -centroids of R and of the solid of revolution. Prove $x_1 < x_2$.

Formal statement does not define R , centroid, or solid of revolution – it uses the analytic formulas directly.

```
theorem putnam_2025_b2
  (f : ℝ → ℝ)
  (hf_cont : ContinuousOn f (Icc 0 1))
  (hf_mono : StrictMonoOn f (Icc 0 1))
  (hf_nonneg : ∀ x ∈ Icc (0 : ℝ) 1, 0 ≤ f x) :
  (∫ x in (0:ℝ)..1, x * f x) /
    (∫ x in (0:ℝ)..1, f x) <
  (∫ x in (0:ℝ)..1, x * (f x) ^ 2) /
    (∫ x in (0:ℝ)..1, (f x) ^ 2) := by
  sorry
```

Geometric narrative replaced by analytically equivalent integral ratios. The π cancels; denominator positivity is a proof obligation, not a hypothesis.

Gemini Draft: Correct Shape, Lean Syntax Fixes

Two layers on this slide

1. Apply Lean syntax fixes from earlier sessions:

```
import Mathlib, open Set Real
```

```
MeasureTheory Interval, noncomputable  
section, Real.pi, parameterize over f.
```

2. Once it typechecks, Gemini's **mathematical shape** is right.

```
noncomputable section  
def area_R (f : ℝ → ℝ) : ℝ := ∫ x in (0:ℝ)..1, f x  
def x1 (f : ℝ → ℝ) : ℝ := (∫ x in (0:ℝ)..1, x * f x) / area_R f  
def vol_solid (f : ℝ → ℝ) : ℝ := Real.pi * ∫ x in (0:ℝ)..1, (f x)^2  
def x2 (f : ℝ → ℝ) : ℝ :=  
  (Real.pi * ∫ x in (0:ℝ)..1, x * (f x)^2) / vol_solid f
```

Both versions express the same analytic comparison: AxiomMath inlines the integrals; Gemini names them. Presentation choice, not error.

"Indefinite Integral" Is Not A Defined Object

(A) Derivative side

```
HasDerivAt : ( $\mathbb{k} \rightarrow F$ )  $\rightarrow$   $F \rightarrow \mathbb{k} \rightarrow \text{Prop}$   
-- predicate:  $F' = f \ x \text{ at } x$   
deriv : ( $\mathbb{k} \rightarrow F$ )  $\rightarrow \mathbb{k} \rightarrow F$   
-- total function; junk value 0 if no deriv
```

In

`Mathlib.Analysis.Calculus.Deriv.Basic`.
`HasDerivAt` is a `Prop`, `deriv` is a value.

(B) Definite integral side

```
intervalIntegral  
  : ( $\mathbb{R} \rightarrow E$ )  $\rightarrow \mathbb{R} \rightarrow \mathbb{R} \rightarrow \text{Measure } \mathbb{R} \rightarrow E$   
-- notation:  $\int x \text{ in } a..b, f \ x \ d\mu$ 
```

In

`Mathlib.MeasureTheory.Integral.IntervalIntegral`.
built on Bochner `MeasureTheory.integral`.

There is **no** `indefiniteIntegral` / `antiderivative` / `primitive` def in `mathlib`. Any "indefinite integral" must reduce to (A) or (B); FTC is the theorem connecting them, not a definition.

Three Informal Phrases, Two Primitives

Natural language	Real identity	Mathlib expression
" F is an antiderivative of f "	predicate on side (A)	$\forall x, \text{HasDerivAt } F (f\ x) x$
"the antiderivatives of f "	set carved out by that predicate	$\{F \mid \forall x, \text{HasDerivAt } F (f\ x) x\}$
" $\int_a^x f(t) dt$ "	a function on side (B)	$\text{fun } x \Rightarrow \int t \text{ in } a..x, f\ t$

The first row uses **only** the derivative side. There is no integral in it.

The second row is a **set-builder**, not a `def`. mathlib has no `antiderivatives` definition.

The third row uses **only** the integral side. It is just a lambda over `intervalIntegral`.

Every alignment of "the indefinite integral of f " has to commit to one of three forms above – and therefore to one of the two primitives (A) or (B).

`deriv f x = 0` Is **Not** Evidence Of Differentiability

`deriv f x = 0` does **not** imply f is differentiable at x .

Why

`deriv` is **total**. If f not differentiable at x , `deriv f x` returns `0` by convention (junk value).

Two interpretations

1. f is differentiable, derivative happens to be 0;
or
2. f not differentiable, `deriv` returned default 0.

Can't tell from the equation alone.

Correct shape for " $f'(0) = 0$ ": use `HasDerivAt f 0 0` (packs existence + value), or hold `DifferentiableAt ℝ f 0` as a hypothesis before reading `deriv`.

Heaviside Step: The Junk-Value In Action

Setup

$$H(x) = \mathbf{1}_{[0,\infty)}(x)$$

```
def H : ℝ → ℝ :=  
  Set.indicator (Set.Ici 0) (fun _ => 1)
```

H is **not continuous** at 0, so certainly not differentiable there.

Yet in Lean

```
example : deriv H 0 = 0 := by sorry
```

This holds – but via the **junk-value** path, not because the derivative equals 0.

```
deriv H 0 = 0 ✓ (true)  
DifferentiableAt ℝ H 0 × (false)  
HasDerivAt H 0 0 × (false)
```

Reading `deriv f x = 0` as "the derivative at x is 0" is wrong whenever the function is not differentiable there. The compiler accepts the equation; the mathematical content has silently dropped out.

Fundamental Theorem Of Calculus As Statement Shapes

FTC-1

If $F(x) = \int_a^x f(t) dt$, then $F'(x) = f(x)$.

```
HasDerivAt  
(fun x => ∫ t in a..x, f t) (f x) x
```

FTC-2

If $F'(x) = f(x)$, then $\int_a^b f(x) dx = F(b) - F(a)$.

```
(∫ x in a..b, f' x) = F b - F a
```

Mathlib theorem names:

```
intervalIntegral.integral_eq_sub_of_hasDerivAt  
intervalIntegral.integral_deriv_eq_sub'
```

The real content lives in the statement

- `HasDerivAt`, `HasDerivWithinAt`, or `deriv`?
- Where does the derivative hold?
- Integrability hypothesis?
- Oriented interval?
- Codomain structure?

"The fundamental theorem of calculus" $\implies$ a family of theorem schemas.

Worked Example: $f' \equiv 0 \implies f$ Constant (1/2)

Informal. f continuous on $[a, b]$, $f' = 0$ on $(a, b) \Rightarrow f$ constant on $[a, b]$.

"Obvious" formal shape

```
theorem foo {a b : ℝ} (hab : a < b) {f : ℝ → ℝ}
  (hcont : ContinuousOn f (Icc a b))
  (hdiff : DifferentiableOn ℝ f (Ioo a b))
  (hderiv : ∀ x ∈ Ioo a b, deriv f x = 0) :
  ∀ x y ∈ Icc a b, f x = f y := sorry
```

Reads off the informal sentence directly.

No Mathlib lemma matches

Closest in

`Mathlib.Analysis.Calculus.MeanValue:`

- `constant_of_derivWithin_zero`
- `constant_of_has_deriv_right_zero`

None take `DifferentiableOn + deriv = 0` on `Ioo a b`. The gap: at endpoints Mathlib wants `derivWithin / HasDerivWithinAt`, not `deriv`.

Worked Example: $f' \equiv 0 \implies f$ Constant (2/2)

Real Mathlib API

```
-- closed interval  $\impl$  derivWithin, not deriv
example
  {a b : ℝ} {f : ℝ → ℝ}
  (hdiff : DifferentiableOn ℝ f (Icc a b))
  (hderiv : ∀ x ∈ Ico a b,
    derivWithin f (Icc a b) x = 0) :
    ∀ x ∈ Icc a b, f x = f a :=
  constant_of_derivWithin_zero
  hdiff hderiv
```

The lemma is `constant_of_derivWithin_zero`.
Same lesson as the LeanArchitect / Multivariate Taylor case at the end of Part IV: at endpoints, use the `Within` variants.

Four alignment shifts vs. draft

- `DifferentiableOn` on `Icc` (continuity inherited)
- `derivWithin`, not `deriv`
- condition on `Ico a b`, not `Ioo a b`
- conclusion pinned to $f(a)$, not symmetric

The "deriv = 0 $\impl$ constant" theorem has at least two real Mathlib shapes – both use the `Within` variants. **Neither** uses `deriv` or `HasDerivAt` directly.

Mathlib Does Not Start From Textbook Notation

Notation used by the Putnam B2 statement earlier:

```
∫ x in a..b, f x
```

comes from

`MeasureTheory.Integral.IntervalIntegral`. Not Riemann sums as primary definition; an oriented interval integral built on the measure-theoretic integral.

```
∫ x in a..b, f x ∂μ  
= ∫ x in Ioc a b, f x ∂μ  
- ∫ x in Ioc b a, f x ∂μ
```

Orientation is part of the API:

```
∫ x in a..a, f x ∂μ = 0  
∫ x in b..a, f x ∂μ  
= - ∫ x in a..b, f x ∂μ  
∫ x in a..b, f x ∂μ  
+ ∫ x in b..c, f x ∂μ  
= ∫ x in a..c, f x ∂μ
```

Formal libraries often choose definitions for theorem engineering, not for matching first-semester textbook wording.

Bochner Integral: Lebesgue For Vector-Valued Functions

Concrete example: velocity $\rightarrow$ displacement

$\mathbf{v} : [0, T] \rightarrow \mathbb{R}^3$ – particle velocity in 3-space.

$$\int_0^T \mathbf{v}(t) dt \in \mathbb{R}^3$$

A **vector** integral. Integrate each coordinate Lebesgue-style; collect into one vector.

Bochner: same idea, any value type

$f : \mathbb{R} \rightarrow E, \int f dx \in E$. E = any space with "add" and "measure size" ($\mathbb{R}, \mathbb{R}^n, \mathbb{C}, \dots$).

$E = \mathbb{R} \Rightarrow$ Bochner = Lebesgue. Nothing new for this course.

When Mathlib says `MeasureTheory.integral`, read "Lebesgue integral". The Bochner name signals the library kept the door open for vector / complex / matrix cases – one definition proved once, specialises to all.

The Underlying Integral Is Bochner Integral

`MeasureTheory.integral` is the Bochner integral:

```
input: measure  $\mu$ ,  $f : \alpha \rightarrow E$ , structure on  $E$   
output: element of  $E$ 
```

Not only real-valued – takes values in any suitable normed vector space.

Notations:

```
 $\int x, f \, x \, d\mu$   
 $\int x \text{ in } s, f \, x \, d\mu$   
 $\int x \text{ in } a..b, f \, x \, d\mu$ 
```

Hierarchy: measurable space / measure $\rightarrow$
integrable functions $\rightarrow$ Bochner integral $\rightarrow$ set
integral $\rightarrow$ interval integral.

Mathlib Also Has Riemann-Style Integrals

Riemann-style lives in `BoxIntegral`:

```
BoxIntegral.HasIntegral  
BoxIntegral.Integrable  
BoxIntegral.integral  
BoxIntegral.IntegrationParams.Riemann
```

Same framework hosts Henstock-Kurzweil and McShane via different `IntegrationParams`.

Textbook phrase → Mathlib object:

```
"the Riemann integral"  
  → BoxIntegral + IntegrationParams.Riemann  
"integral from a to b"  
  → intervalIntegral over MeasureTheory.integral
```

Bridge theorems connect the two under suitable hypotheses. Multiple definitions, one common proof interface.

What The Integral Case Teaches About Alignment

Template for the rest of the lecture:

informal phrase

"the integral of f from a to b "

formal choices

type of f

codomain structure

measure

interval convention

integrability condition

endpoint behavior

intended theorem shape

An informal mathematical statement is often a compressed reference to a whole formal design space.

The integral example was concrete. Part II abstracts it into a general checklist.

Part II. General Mechanism

From informal sentence to formal statement

The Statement Is The Object

Week 8 builds:

a formal theorem statement

(not a proof)

Two checks on a formal statement:

1. It typechecks in the formal language.
2. It expresses the intended mathematical meaning of the informal statement.

Dangerous autoformalization failures pass (1) but fail (2).

Example:

Informal: *Every subgroup of a cyclic group is cyclic.*

A type-correct but wrong formalization could:

- state the ambient group is cyclic in the conclusion instead of the subgroup;
- use the wrong formal definition of "cyclic".

Alignment Checklist

1. **Objects.** Variables, functions, sets, structures, parameters.
2. **Types.** Each object's type.
3. **Structures.** Algebraic, order, topological, metric, measure, category.
4. **Quantifiers.** Universal, existential, order.
5. **Assumptions.** Explicit vs convention-hidden.
6. **Definitions.** Which library definition corresponds.
7. **Notation.** Overloaded symbols disambiguated.
8. **Statement strength.** Too weak, too strong, just right.
9. **Edge cases.** 0, empty set, equal endpoints, boundaries.
10. **Back-translation.** Translate the formal back; does it still read as the original theorem?

This is the human version of statement validation. Modern systems (Part IV) automate parts of it.

Typechecking Is Not Semantic Alignment

A Lean statement can be well-typed and still be the wrong theorem.

Common failure modes

- **Too weak.** Only a special case.
- **Too strong.** Extra hypotheses, or sharper conclusion.
- **Wrong object.** Nearby but different definition.
- **Wrong scope.** Quantifier order swapped.

- **Wrong ambient structure.** Specialized to $\mathbb{R}$ instead of generic.
- **Wrong convention.** Index from 0 vs 1; open vs closed interval; left derivative vs derivative.

Formal statement checking is not enough. We also need statement alignment.

Part III. Worked Examples

Simple to research-level alignment problems

The First n Odd Numbers Sum To n^2

Informal:

The sum of the first n odd numbers is n^2 .

Hidden choices

- n is a natural number?
- "first n " means n terms, or up to n ?
- Odd numbers as $2i + 1$ or $2i - 1$?
- Index $0..n - 1$ or $1..n$?
- Sum over `Finset.range`, list, recursion?
- Equality in `Nat`, `Int`, `Real`?

A possible Lean statement:

```
example (n : Nat) :  
  (∑ i in Finset.range n, (2 * i + 1))  
    = n ^ 2 := by  
  ...
```

This statement has already made these choices:

- `n : Nat`;
- zero-based indexing;
- first n terms are $1, 3, \dots, 2n - 1$;
- finite sum over `Finset.range`;
- equality in `Nat`.

Set Distributivity

Informal:

$$A \cap (B \cup C) = (A \cap B) \cup (A \cap C)$$

Hidden choices

- Universe type α ?
- A, B, C are `Set  $\alpha$` ?
- Extensional set equality?
- $\cap, \cup$ as set ops or lattice ops?

```
variable { $\alpha$  : Type*} (A B C : Set  $\alpha$ )
```

```
example :
```

```
  A  $\cap$  (B  $\cup$  C)  
    = (A  $\cap$  B)  $\cup$  (A  $\cap$  C) := by  
  ...
```

Even this simple example forces

- a universe α
- objects typed as `Set  $\alpha$`
- set lattice notation
- extensional equality

The Inverse Of A Product

Informal:

The inverse of a product is the product of the inverses in reverse order.

Hidden choices

- Group, monoid-with-inverses, division ring, or category?
- Multiplication commutative or not?
- Conclusion $(ab)^{-1} = b^{-1}a^{-1}$?
- In a commutative setting, reverse order is hidden; otherwise substantive.

```
variable {G : Type*} [Group G]

example (a b : G) :
  (a * b)-1 = b-1 * a-1 := by
  ...
```

Natural language says "product" without naming the algebraic structure. Formal language cannot use multiplication until the structure is fixed.

Every Subgroup Of A Cyclic Group Is Cyclic

Informal:

Every subgroup of a cyclic group is cyclic.

Hidden choices

- Ambient group G .
- Multiplicative vs additive notation.
- "Cyclic" as a predicate on groups, subgroups, or modules?
- Subgroup as bundled `Subgroup G`?
- Subgroup cyclicity uses the inherited group structure?

```
variable {G : Type*} [Group G]

-- sketch shape; final mathlib API
-- packages this differently:
∀ H : Subgroup G,
  IsCyclic G → IsCyclic H
```

"Subgroup" in natural language refers both to a subset-like object and to a group in its own right. Formalization must choose the bundled representation that exposes the needed structure.

Continuous Functions On $[a, b]$ Attain A Maximum

Informal:

A continuous function on $[a, b]$ attains a maximum.

Hidden choices

- $a \leq b$ assumed?
- $f : \mathbb{R} \rightarrow \mathbb{R}$, or $f : X \rightarrow \mathbb{R}$?
- `Continuous f` vs `ContinuousOn f (Set.Icc a b)`?
- "Attains a maximum" means $\exists x \in [a, b], \forall y \in [a, b], f(y) \leq f(x)$?
- Is this really compact-set extreme value?

$$\begin{aligned} \exists x \in \text{Set.Icc } a \ b, \\ \forall y \in \text{Set.Icc } a \ b, \ f \ y \leq f \ x \end{aligned}$$

A formal statement here specifies:

- domain set: `Set.Icc a b`;
- codomain order: $\leq$ on $\mathbb{R}$;
- continuity predicate: `ContinuousOn`;
- compactness source: closed bounded interval.

A common theorem name often hides compactness, order, and domain restriction.

$\text{Gal}(E/\mathbb{Q}) \cong Q_8$: One Statement, Several Formal Shapes

Informal:

Let $\alpha = \sqrt{(2 + \sqrt{2})(3 + \sqrt{3})}$ and $E = \mathbb{Q}(\alpha)$.
Show $\text{Gal}(E/\mathbb{Q}) \cong Q_8$.

Hidden choices

- α as `Real.sqrt ...`;
- E as `IntermediateField.adjoin  $\mathbb{Q}$  { $\alpha$ }`;
- $\text{Gal}(E/\mathbb{Q})$ as `$\mathfrak{r}E \approx_a[\mathbb{Q}] \mathfrak{r}E$` ;
- Q_8 as `QuaternionGroup 2`;
- $\cong$ as `$\approx^*$` (Type, not Prop).

```
-- Shape A: explicit MulEquiv (def, not theorem)
def gal_iso_Q8 :
  ( $\mathfrak{r}E \approx_a[\mathbb{Q}] \mathfrak{r}E$ )  $\approx^*$  QuaternionGroup 2 := by
  sorry

-- Shape B: Galois-group predicate
theorem q8_is_galois_group
  [MulSemiringAction (QuaternionGroup 2)  $\mathfrak{r}E$ ] :
  IsGaloisGroup
    (QuaternionGroup 2)  $\mathbb{Q}$   $\mathfrak{r}E$  := by
  sorry

-- Shape C: finite Galois group object
noncomputable def q8_as_finGaloisGroup_iso :
  Efin.finGaloisGroup
     $\equiv$  FiniteGrp.of (QuaternionGroup 2) := by
  sorry
```

Three Shapes, Three API Paths

Shape A – MulEquiv

Direct group isomorphism.

$(\mathbb{1}E \simeq_a [\mathbb{Q}] \mathbb{1}E) \simeq^* \text{QuaternionGroup } 2$

Closest to the informal $\cong$. Must use `def` (the type is `Type`, not `Prop`).

Shape B – IsGaloisGroup

Q_8 acts faithfully on E with fixed field $\mathbb{Q}$.

`IsGaloisGroup (QuaternionGroup 2)  $\mathbb{Q}$   $\mathbb{1}E$`

Mathlib bridges to `MulEquiv` via `IsGaloisGroup.mulEquivAlgEquiv`.

Shape C – FiniteGrp

Package E as `FiniteGaloisIntermediateField`.

`Efin.finGaloisGroup  $\cong$  FiniteGrp.of (QuaternionGroup 2)`

"Finite Galois" lives inside the object, not as a separate hypothesis.

Three reasonable formal targets for one informal isomorphism. Choice of shape changes available API and downstream proof routes.

Does $\sum_{n=1}^{\infty} \frac{1}{n}$ Converge?

Hidden commitments

NL	hidden choice
$\sum_{n=1}^{\infty}$	partial-sum vs limit object
$\frac{1}{n}$	$\mathbb{Q}$ -, $\mathbb{R}$ -, $\mathbb{C}$ -valued?
"converge"	classical limit, or $= +\infty$?

Three Lean shapes (equivalent in $\mathbb{R}$)

```
∃ L, Tendsto (fun N => ∑ k in range N, 1/(k+1)) atTop (ℕ L)
Summable (fun n : ℕ => (1:ℝ)/(n+1))
HasSum (fun n : ℕ => (1:ℝ)/(n+1)) L
```

Answer (Euler, 1734)

$\sum 1/n = +\infty$ via integral comparison:
 $\int_1^{N+1} dx/x \leq \sum_{n=1}^N 1/n$, LHS = $\ln(N+1) \rightarrow \infty$.

Alternating: opposite story

$\sum (-1)^{n-1}/n = \ln 2$, converges **conditionally**.
Lebesgue/Bochner refuses it – $\sum |a_n|$ must be finite.

Trap: $\neg \text{Summable } (1/n) \checkmark$ and $\text{tsum } (1/n) = \top \checkmark$ – but the second uses $\mathbb{R}_{\geq 0\infty}$, not $\mathbb{R}$. Pattern-matching "diverges" $\rightarrow "= \infty"$ picks wrong target.

Part IV. Modern Autoformalization Systems

Mechanisms aimed at the alignment problem

What Modern Systems Are Trying To Fix

Five failure shapes

- **Free-variable swap** – `pi` $\neq$ `Real.pi`
- **Default-argument trap** – `(f := ...)` not asserted
- **Wrong object** – `Continuous` vs `ContinuousOn`
`s`
- **Junk-value** – `deriv f x = 0` when not diff'ble
- **Missing hypothesis** – silent drop of `[Fintype G]`

The formal statement **compiles** and is **not** the intended theorem.

Each paper attacks one piece

Cog	Paper
Stronger acceptance	Rethinking – BEq
Library grounding	Rethinking, Aria, LoC
Self-critique	ReForm
No gold reference	Roundtrip
FL $\rightarrow$ NL infra	Herald / LeanSearch
Project scale	LeanArchitect

Rethinking: Typecheck Is Too Weak

*Lean accepts the statement
does **not** imply
the statement is the intended theorem.*

The fix: BEq

Bidirectional Extended definitional Equivalence:

both directions $S_{\text{gen}} \leftrightarrow S_{\text{ref}}$ must derive under a bounded set of rewrites.

Single direction can silently weaken the claim.

Both directions cannot.

On 200 expert-labeled pairs

Metric	Prec.	Acc.
Typecheck	35	35
BLEU	63	69
LLM judge	71	83
BEq	100	91

BEq's recall is 73 % – the precision-100 zone. The gap above 73 % is where Aria / ReForm / Roundtrip do their work.

Rethinking: Retrieval Beats Generation

Informalized Mathlib

243 k objects + 139 k theorems → bottom-up
informalisation gives each object NL description +
dependency list.

Old prompt: *"Translate this to Lean."*

RAutoformalizer prompt: *"Available:*

*Complex.differentiableOn, ... Now write the
statement."*

Con-NF BEq@8

baseline:	4.6 %	
+ retrieval:	16.9 %	(+12 pp)
+ oracle:	~35 %	(+18 pp)

Plain semantic search misses Recall@5 $\approx$ 35 %
(ProofNet), 24 % (Con-NF).

Bottleneck is not Lean-writing skill – it is which Mathlib objects sit in the context. Prompt should be "identify Mathlib objects for each concept, then compose"; Aria / LoC / ReForm operationalise this.

Aria: Concept Dependency Graph Before Lean Code

Graph-of-Thought pipeline

```
NL → concept graph
    → ground leaves (LeanSearch)
    → bottom-up synthesis
    → compiler-in-the-loop
    → Lean theorem
```

Never write Lean from NL directly.

"First n odd numbers sum to n^2 "

```
      target
      |
sum / first-n-odd / n²
|      |      |
Finset Finset. literal
.sum   range +
      (2k+1)
```

Every leaf grounds in Mathlib; inner nodes synthesized.

Headlines: ProofNet 68.5 %, FATE-H 71 %, homological 42.9 %. Ablations: remove RAG → 0 % on conjectures; remove GoT → 6/14 → 1/14. Both retrieval and graph structure load-bearing.

AriaScorer: Term-Level Grounding

Question

A generated definition **compiles** – how do we know it is the intended **mathematical concept**?

Mechanism

For each term in the candidate Lean, `jixia` retrieves from informalized Mathlib: name, type, def body, informal name + description. A judge compares assumptions/conclusions in this *grounded* context, not on surface strings.

Traps surface checkers miss

UFD / IsDomain. Surface checker flags missing `IsDomain`; AriaScorer sees `UniqueFactorizationMonoid` already inherits it.

QuaternionGroup 1 vs 2. `QuaternionGroup n` has order $4n$. Surface match accepts either; AriaScorer reads the index – wrong for Q_8 .

Makes Aria **self-correcting**: a compiled-but-wrong synthesis triggers another GoT pass with the failed term marked for re-grounding.

LoC-Decomp: A Template That Forces Commitments

The seven sections

1. imports
2. aux types
3. mathematical systems
4. aux functions
5. constraints
6. solution constraints
7. final statement

Each forces a commitment NL hides.

Worked: "find all n with n^2+n+1 prime"

One-shot LLM:

```
theorem foo (n : ℕ) :  
  Nat.Prime (n2+n+1) ↔ n = 1
```

Wrong on 3 dimensions. Template forces:

```
answer : Set ℕ = {n | n ≥ 1 ∧ Nat.Prime (n2+n+1)}  
theorem foo : answer = {1}
```

Three commitments NL fuses get separated: *what the unknowns are* / *what the answer is* / *what "find" means*.
Same 7 sections work for "find X ", "show P ", inequality, existence.

LoC-Decomp: ASCC + Numbers

Four-cell diagnostic

Compile	ASCC	Diagnosis
✓	✓	alignment evidence
✓	×	wrong theorem
×	✓	math right, API wrong
×	×	broken at both

ASCC back-translates **per template section**;
discrepancies localise to a Lean line.

DeepSeek-V3 + LoC pass rate


	MATH-500	miniF2F
raw LLM	~35 %	—
+ template	63 %	77 %
+ ASCC loop	84 %	90 %

ASCC checker itself (150 expert pairs): P 0.90, R 0.82.

Loop does as much work as template: structured output alone leaves 20 pp; per-segment feedback closes it. ReForm next moves the loop *inside* the model.

ReForm: Reflection Loop + PBSO Reward

Architecture



$\mathcal{L} = -\mathbb{E}[R_{\text{task}} + \lambda \sum_t R_{\text{aux}}(t)]$. Without R_{aux} the model rubber-stamps S_1 .

Shallow vs specific critique

Shallow: "Syntax valid; matches problem." $\rightarrow R_{\text{aux}} \approx 0.1$

Specific: "Q is conditional 'if $F' = f$ then ...'. S_1 never introduces F . Missing `HasDerivAt`; integrability." $\rightarrow R_{\text{aux}} \approx 0.85$

A useful critique is not "yes/no on S_t "; it is a **repair plan for S_{t+1}** .

ReForm: Numbers + Audit Of "Gold" Benchmarks

Headline (semantic-consistency %)

Model	miniF2F	ProofNet	Putnam	AIME
Baseline best	73	48	36	22
ReForm-32B	91	70	62	67
Δ	+18	+23	+27	+45

Pattern: **harder benchmark, larger lift**. Average +22.6 pp.

Audit of "gold" benchmarks

ReForm's ConsistencyCheck (859 items) on **human-written** gold:

miniF2F: **16.4 %** semantic errors. ProofNet: **38.5 %**.

Caught: $\sqrt{11} \rightarrow 11$; $\deg p \leq 80 \rightarrow \deg p < 80$;
"find X " baked into " $X = c$ ".

Gold references carry 16-39 % errors $\Rightarrow$ "compare to reference" is not clean ground truth. **Roundtrip** (next) drops the gold assumption.

Roundtrip: Verification Without A Gold Target

Roundtrip

```
y_orig = formalize(x)      T1
x'      = informalize(y_orig) T2
y_rt    = formalize(x')    T3

check  y_orig ≡ y_rt
```

Disagreement $\Rightarrow$ drift in T1, T2, or T3.

Scoped repair + debug

Broke	Repair
T1	re-formalize, re-run T2/T3
T2	only re-run T3
T3	re-formalize y_{rt} only

Debug Qs: was T1 wrong? was T2 misleading?
was T3 unstable?

Parallels LoC-Decomp: both **localise** the failure (segment vs stage). The only Part IV mechanism without any gold reference.

Roundtrip: Numbers + Wrong Fixed Points

Statute accuracy (Z3 equivalence)

Texas Transportation / Parks & Wildlife codes:

No repair:	45 / 66 %
Claude + scoped repair:	85 / 86 %
GPT + Claude diagnoser:	91 / 95 %

GPT's default prompt blames T1 in **83 / 97 %** of cases; a non-sequential prompt drops to **9 / 16 %**. Same models, only the diagnoser changed.

Wrong fixed point

T2 drops the same hypothesis as T1 $\Rightarrow$ roundtrip stabilises on wrong $y_{\text{orig}} \equiv y_{\text{rt}}$. Mitigation: multiple T2 paraphrase paths.

Why statute, not Lean

Z3 decides FOL equivalence; Lean equivalence needs Mathlib lemmas – same idea, harder checker.

Why The Reverse Direction Matters: Formal $\rightarrow$ Informal

Four uses for FL $\rightarrow$ NL

1. Training data. Mathlib4 has ~200 K formal decls, ~0 paired NL. Herald produces 44 k pairs – all by back-translation.

2. Retrieval. LeanSearch, Aria, ReProver index *informalized* Mathlib, not raw source.

3. Verification. ReForm, Roundtrip, Herald validator all back-translate $S \rightarrow Q'$ and compare $Q \approx Q'$. No FL $\rightarrow$ NL $\rightarrow$ no check.

4. Documentation. 1.5 M lines of Mathlib become navigable. Same component powers docs, errors, tutorials.

FL $\rightarrow$ NL is **not symmetric** with NL $\rightarrow$ FL. It is the prerequisite that makes the rest of Part IV possible.

Herald: Mathlib4 → Natural-Language Dataset

Pipeline

Mathlib4 → jixia → informalize → Herald dataset + Translator

Scale: 580 k statements, 44 k NL↔FL theorem pairs, 45 k proof pairs.

Same LLM, same statement – the difference is **structured metadata fed to the prompt**, not model capacity.

Set.Ico walkthrough

```
theorem foo (a b c : ℝ) :  
  c ∈ Set.Ico a b ↔ a ≤ c ∧ c < b
```

Naive: "c belongs to Ico from a to b iff $a \leq c$ and $c < b$." – identifier leaks.

Herald: "c lies in the left-closed right-open interval $[a, b)$ iff $a \leq c < b$."

LeanSearch: Informalized Mathlib As A Service

Pipeline

```
NL query → embedding (Herald pairs)
          → nn over informalized Mathlib
          → Mathlib declarations
```

Example: "*continuous function on compact set attains max*" → `IsCompact.exists_isMaxOn`, `ContinuousOn.isCompact_image`. Deployed at leansearch.net.

Downstream

Consumer	Use
LeanSearch	NL → API search
Loogle	type-pattern + NL labels
Aria / ReProver	retrieval in NL → FL

FL → NL is the **upstream infrastructure** behind every "RAG over Mathlib" pipeline in Part IV – not a symmetric counterpart of NL → FL.

Herald Translator: Numbers + Two Case Studies

Coverage

miniF2F-test	96.7 %
internal graduate	23.5 %
College CoT	16.0 %

Pipeline: compile $\rightarrow$ back-translate $\rightarrow$ NLI.

Case 1: quantifier order

```
 $\forall x, \forall \epsilon, \exists \delta, \dots$  -- continuity  
 $\forall \epsilon, \exists \delta, \forall x, \dots$  -- uniform
```

Moving x past $\exists \delta$ rewrites the theorem.

Case 2: Four Lemma (mono half)

Lean header full of `ComposableArrows.app'  $\varphi$  i`
... index proofs. Faithful NL: "*morphism between exact chains; if φ_0 epi, φ_1, φ_3 mono, then φ_2 mono.*"

Faithful FL $\rightarrow$ NL must restore the **mathematical surface**, not Lean implementation details.

What Makes Formal $\rightarrow$ Informal Hard: Three Tax

1. Naming tax

Mathlib identifiers = encyclopaedic labels.

```
Set.Ico  $\rightarrow$  [a,b)  
StrictMono  $\rightarrow$  strict inc.  
IsCompact  $\rightarrow$  compact  
Mono/Epi  $\rightarrow$  mono/epi
```

Need a learned naming map.

2. Type-class tax

Which implicit args surface?

- `[Abelian C]` $\rightarrow$ yes
 - `[DecidableEq G]` $\rightarrow$ no
 - `[CommRing R]` $\rightarrow$ yes
 - universes $\rightarrow$ no
- Heuristic, per Mathlib design.

3. Surface tax

- chain $a \leq c \wedge c < b \rightarrow a \leq c < b$
- bracket notation
- `mono_of_...` $\rightarrow$ "Four Lemma"
- $\mathbb{R}$ vs `Real`

Three taxes independent. Naming alone leaves `DecidableEq` noise; type-class alone leaves un-chained algebra; surface alone leaves leaking identifiers.

Beyond Theorem Headers: Three Levels Of FL → NL

Level 1: statement

```
ContinuousOn f s → IsCompact s  
  → IsCompact (f '' s)
```

→ "Continuous on compact
⇒ compact image."

Herald: 97 % miniF2F, 16 %
graduate.

Level 2: proof

```
intro ε hε  
obtain ⟨δ, hδ, h⟩ := ...  
calc |f y - f x| < ε/2 ...
```

→ NL prose with proof-state
context at every step.

Lean-STaR partial; mostly
open.

Level 3: definition

```
class CompleteLattice α extends ...  
  sSup : Set α → α  
  le_sSup ...
```

→ "Lattice with arbitrary
suprema."

Mathlib docstrings, hand-
written; no scalable
generation.

Most Part IV "autoformalization needs autoinformalization" claims assume **Level 1**. Levels 2 and 3 reopen the alignment problem in harder forms.

LeanArchitect: @[blueprint] – One Source Of Truth

The Lean attribute

```
@[blueprint "thm:add-comm"  
  (statement := /-- Addition commutative. -/)  
  (uses := [def:nat, lem:zero-add]))  
theorem MyNat.add_comm ...
```

Holds label + NL + deps. Extractor reads actual used constants and emits `\leanok`, `\uses{ ... }` automatically.

Drift modes auto-sync kills

- stale prose vs Lean header
- missing/wrong `\uses{ ... }`
- wrong `\leanok`
- divergent edits

LeanArchitect makes Lean the single source; LaTeX blueprint is *derived*.

Caveat. Auto-sync $\neq$ semantic equivalence. @[blueprint] closes *metadata* drift; whether NL and Lean statements mean the same thing is the alignment problem Aria / LoC / ReForm address.

LeanArchitect Case: Multivariate Taylor And `derivWithin`

Workflow

GPT-5 drafts blueprint + sorrys → LeanArchitect graph → Aristotle proves leaves → human inspects failures.

What happened

~40 nodes; **37 discharged**; **3 stuck**. All on derivatives *on the segment* $[a, b]$.

The bug – not a proof problem

Lean used `deriv` for a function on segment $[a, b]$.
Correct: `derivWithin (Icc a b)`. Outside the segment `deriv` returns junk; the statement was never the segment-restricted theorem.

With the graph: trace upstream → fix the statement, not the prover.

Same alignment trap, three scales – Part I's `deriv` junk value, Part III's `ContinuousOn` vs `Continuous`, here `deriv` vs `derivWithin` at project scale.

Part IV Synthesis

Better acceptance

What the statement *means*, not just compiles.

- BEq (precision 1.00)
- AriaScorer term grounding
- ConsistencyCheck audit

Library grounding

Autoformalization = retrieval + generation.

- RAutoformalizer (+12 pp)
- Aria GoT pipeline
- Herald + LeanSearch (infra)

Feedback + scale

Alignment is a process.

- LoC ASCC loop (+20 pp)
- ReForm + PBSO (+22.6 pp)
- Roundtrip (no gold)
- @[blueprint] at scale

Seven papers, one diagnosis: **typecheck is the start line, not the finish line**. The Part V activity is a hand-run miniature of BEq + AriaScorer + ASCC + Roundtrip.

Part V. Alignment Template

Fields a formal statement must pin down

Fields A Formal Statement Must Pin Down

Informal statement:

Objects:

Types:

Structures / typeclasses:

Definitions from library:

Quantifiers:

Assumptions:

Conclusion:

Edge cases:

Back-translation:

The deliverable is not a proof. It is a formal statement whose every field above has been committed to deliberately.

Sample Informal Claims For Alignment

A. The inverse of a product is the product of inverses in reverse order.

B. A continuous function on $[a, b]$ attains a maximum.

C. Every subgroup of a cyclic group is cyclic.

D. The first n odd numbers sum to n^2 .

E. The indefinite integral of f is F .

Three possible formal targets: derivative-side predicate, derivative-side set, or a specific definite-integral-with-variable-upper-limit. Each gives a different theorem.

Alignment Audit: Six Questions For Any Formal Statement

Six questions to check a formal statement S against an informal claim Q :

1. If S is true, does Q follow?
2. If Q is true, does S follow?
3. Was an assumption silently added that wasn't there in Q ?

4. Was a needed assumption dropped?
5. Was the object swapped (e.g. ambient group instead of subgroup, `deriv` instead of `derivWithin`)?
6. Does a back-translation of S still sound like Q ?

"Yes" on (1) and (2), "no" on (3)-(5): the minimum bar for semantic alignment. Typecheck does not give this for free.

Part VI. Analytic Number Theory

Harmonic series, primes, and density

Leonhard Euler (1707–1783)



Handmann pastel, 1753

Highlights

- Proved $\sum 1/n = \infty$
- Euler product: $\zeta(s) = \prod_p (1 - p^{-s})^{-1}$
- First analytic proof of infinitely many primes (via $\sum 1/p$ diverging)
- Standardised $\sum, e, i, \pi, f(x)$

Euler's identity $\sum_p \ln \frac{1}{1-1/p} = \sum_n 1/n$ is one line that hides type choice, convergence condition, and infinite-product definition. The first alignment question is: *what does this $\sum$ actually mean?*

Harmonic Series: $\sum \frac{1}{n}$ Diverges

Integral comparison

$$\sum_{n=1}^N \frac{1}{n} \geq \int_1^{N+1} \frac{dx}{x} = \ln(N+1) \rightarrow \infty$$

Cauchy's blocking

$1 + \frac{1}{2} + (\frac{1}{3} + \frac{1}{4}) + (\frac{1}{5} + \dots + \frac{1}{8}) + \dots -$
infinitely many blocks each $\geq \frac{1}{2}$.

Numerical taste

N	H_N
10	2.93
100	5.19
1000	7.49
10^9	≈ 21.30

Growth $\sim \ln N$; Euler's $\gamma = \lim(H_N - \ln N) \approx 0.5772$, irrationality open.

Diverges **slowly** – that's exactly what lets the alternating version converge.

Alternating Harmonic Series

Convergence (Leibniz)

a_n positive, decreasing, $a_n \rightarrow 0 \Rightarrow \sum (-1)^{n-1} a_n$ converges. $a_n = 1/n$ qualifies.

Limit

From Taylor of $\ln(1+x)$ at $x=1$ (Abel summation): $\sum (-1)^{n-1}/n = \ln 2$.

Conditional, not absolute

$$\sum |(-1)^{n-1}/n| = \sum 1/n = \infty.$$

Riemann rearrangement: rearrange to any real, $\pm\infty$, or divergent. The value $\ln 2$ depends on order.

Lebesgue connection: Lebesgue/Bochner sum on $\mathbb{N}$ requires $\sum |a_n| < \infty$. Alternating harmonic = canonical case where Riemann accepts a value Lebesgue refuses.

Euler's Proof: Infinitely Many Primes

Euler product (1737)

For $s > 1$, $\zeta(s) = \sum 1/n^s = \prod_p (1 - p^{-s})^{-1}$
by unique factorization.

As $s \rightarrow 1^+$

LHS: $\zeta(s) \rightarrow \sum 1/n = \infty$, so product diverges:
 $\prod_p (1 - 1/p)^{-1} = \infty$.

Take logs


$$\sum_p \ln \frac{1}{1-1/p} = \infty$$

Using $\ln \frac{1}{1-x} = x + x^2/2 + \dots$, the tail
 $\sum_p (1/2p^2 + \dots)$ converges (bounded by
 $\sum 1/n^2$), so:

$$\boxed{\sum_p \frac{1}{p} = \infty}$$

Finitely many primes $\Rightarrow$ finite sum. So infinitely many – and $\sum 1/p$ diverges. Compare $\sum 1/n^2 < \infty$:
primes are "denser than squares".

Bernhard Riemann (1826–1866)



Portrait c. 1863

Highlights

- Extended $\zeta(s)$ to a **meromorphic function** on $\mathbb{C}$ with one simple pole at $s = 1$
- 1859 memoir linked ζ 's zeros to prime distribution
- **Riemann Hypothesis:** zeros on $\Re(s) = 1/2$ – open since
- **Identity theorem:** uniqueness of analytic continuation
- Riemann integral, curvature, mapping theorem

Next slide: " ζ has a pole at $s = 1$ " – one Lean statement invoking the whole continuation +

Meromorphic Functions And The Identity Theorem

"Meromorphic"

$f : U \subseteq \mathbb{C} \rightarrow \mathbb{C}$ holomorphic except at isolated **poles**. Equivalently $f = g/h$ with g, h holomorphic.

Riemann's $\zeta(s)$: meromorphic on $\mathbb{C}$, **single simple pole at $s = 1$** , residue 1.

Identity theorem

Meromorphic f, g on connected open U agreeing on a set with a limit point $\Rightarrow f = g$ on U .

So the meromorphic extension is unique;
 $\zeta(-1) = -1/12$ becomes meaningful only with this.

" ζ extends to a meromorphic function on $\mathbb{C}$ with simple pole at $s = 1$ " hides 5 Lean commitments: complex function / holomorphic subset / order of singularity / uniqueness / residue.

Peter Gustav Lejeune Dirichlet (1805–1859)



Public-domain portrait

Highlights

- **Dirichlet's theorem (1837):** infinitely many primes in every coprime AP
- **Dirichlet characters χ + L-functions $L(s, \chi) = \sum \chi(n)/n^s$** – the proof engine
- **Dirichlet density** – works where natural density doesn't
- Pigeonhole principle, Dirichlet series, BC, function $\chi_{\mathbb{Q}}$

Dirichlet's theorem is the first non-trivial result on prime *distribution*, and the proof launched analytic number theory.

Dirichlet's Theorem And Density

Dirichlet (1837)

For coprime a, q , the AP $a, a + q, a + 2q, \dots$ contains **infinitely many primes**.

Equidistribution

Primes are evenly split among the $\varphi(q)$ coprime residue classes mod q :

$$d_{\text{Dir}}(\{p \equiv a \pmod{q}\}) = \frac{1}{\varphi(q)}$$

Natural density

$$d_{\text{nat}}(S) = \lim_{x \rightarrow \infty} \frac{\#\{p \in S : p \leq x\}}{\pi(x)}$$

Dirichlet density

$$d_{\text{Dir}}(S) = \lim_{s \rightarrow 1^+} \frac{\sum_{p \in S} p^{-s}}{\sum_p p^{-s}}$$

d_{nat} existing $\Rightarrow d_{\text{Dir}}$ equals it. **Converse fails** – picking the wrong formal definition silently changes the theorem.

Why These Three Topics Sit Together
